CIS 550 PROJECT REPORT - MAX SCHEIBER, BRIAN SHI, AND RIGEL SWAVELY

1. INTRODUCTION

We built **EIGHT**, a travel planning website and social network. The name was inspired by an inside joke among CIS 121 TAs and all too much box wine.

2. PROJECT GOALS

We had three goals for this project. One was to build a scalable application at both the database level and the server level. We wanted to design a schema that was easy to build more features on top of, and we also wanted our application to have good performance (i.e. not just be another class project). Next, we wanted to learn more about Node.js. Max had done some consulting work in Node, but it was pretty new to both Brian and Rigel. Finally, we had traditionally used ORMs or ODMs when writing web applications in the past, so we figured that it would be a different and refreshing experience to have to write raw SQL this time around.

3. APPLICATION DETAILS

3.1. Architecture. The project was build as a Node application on the Express framework, and used the Model View Controller architectural pattern. Our models contained the functions to directly interact with the database, including the prepared queries that were run. Our controllers combined our model functions in order to send the information to the views, which displayed the necessary pages to the user.

3.2. Implementation.

3.2.1. *Photos and Videos.* Uploading photos and videos was fairly straightforward. The url referring to the photo or video was given by the user in one of the views, which is then passed into the database to create a new row in its respective table. A corresponding entry is created in Media, which contains the privacy information, and in Owns, which contains the relation to the user that uploaded the photo or video.

Implementing privacy restrictions was done serverside. When loading photos or videos, if the setting is 'private,' it will only display if the current user is the owner. If the setting is 'sharedWithTripMembers,' then it only displays if the current user is a member of the trip that the media is associated with.

3.2.2. *Recommendations.* We used a very simple algorithm to determine recommendations. First, we look at all of the trips that a user is on, and if any of the trips include destinations that are not part of the user's profile, the destination is recommended. Similarly, if there are any members of those trips that the user is not already friends with, then they are recommended as friends.

3.2.3. *Search.* We use a simple search algorithm based on Damerau-Levenshtein distance. This is obviously not the most intelligent search algorithm - we get some weird results toward the bottom of the search page - but it generally works for catching typos and immediatley relevant results. It would have been better to use some sort of PageRank algorithm to group together relevant destinations on a trip, as one possible improvement.

3.2.4. *Caching, NoSQL.* We did our caching in MongoDB, hosted on Amazon in the same region (us-west) as the rest of our stack. We move popular photos to our cache by tracking their view count. After being moved to the cache, we obviously read the binary data from there instead of loading from the URL. We do not currently purge the cache due to the artificial nature of the project, but we could easily implement a least recently used (LRU) strategy if we needed to.

3.2.5. *Additional Extra Credit Features.* We implemented searchable hashtags for classifying media on the website - specifically, photos, trips, and destinations. Unlike standard search, these are evaluated strictly (like Twitter) and are useful for viewing relationships between disparate sets of user-created here.

We also implemented Bing image search integration. When you perform a search on our website, we also show five relevant images from Bing search. This works extremely well.

4. Technical Challenges

Throughout the course of the project, we encountered several technical challenges. The first of such challenges was deploying our web application to the cloud. Although all members of our team had worked on developing web applications in the past, none of us had used platforms similar to Amazon EC2. Figuring out how to set up the account properly with the correct security and firewall settings proved to be a bit of a challenge, but after carefully following the instructions provided through the course we were able to successfully log in to a server. Finally, once we set up the proper software on the server such as Node and Git, we were able to start the Node server and access our application.

In addition, connecting to Oracle from our web application proved to be fairly difficult. Although there are commonly used, well maintained Node packages to connect to most database systems like PostgreSQL, MongoDB, Redis and more, there is not much support for connecting to an Oracle database. Thankfully, we were able to find one package, but installing it turned out to be fairly difficult. The package required software that had to be separately downloaded from Oracle in order to work, and configuring it took several tries both on our local machines and on the server. Once it was finally set up, however, using

the interface to interact with the database was relatively straightforward.

Finally, we had some trouble creating a schema for the database that suited our requirements. We wanted to have one common interface to go between our different media types, such as trips, destinations and photos, and user posted content regarding those media types, such as ratings and comments. Additionally, we wanted to avoid having multiple likes, commments, and ratings tables for each type of media (eg. photo likes, trip likes, etc), and instead have one common set of tables among all of them. However, that would require a relationship from every media table to the ratings, likes, and comments tables, which still seemed wasteful. Additionally, if we were to implement another type of user created content, such as hashtags, it would require changing every table to include a new relationship under this implementation. Instead, we decided to have one central Media table, which had references to every entry in each of our specific media tables. The media table was indexed by a primary key which was a composite of the relevant media's type and its pid in its own table. For example, if we had a photo with primary key 12 in the Photo table, its corresponding primary key in the Media table would be ('Photo', 12). The ratings, likes, and comments tables then simply used the primary key of the entry in the Media table that it referred to, minimizing duplication.

5. Performance Evaluation

We used Apachebench to benchmark our application. We hypothesized that the main bottleneck of our application was the constant calls to Oracle - we assumed we were IO-bound on the server's end, rather than IO-bound at the client's end (i.e. network latency) or CPU-bound. We therefore conducted several tests to determine whether our hypothesis was correct.

5.1. **Local versus EC2.** We first attempted to control for running our application on `localhost` versus hosting it on Amazon. We therefore ran 100 sequential requests to load the profile of username `rigel`, who had a lot of photos and trips.

Description	Mean (sec.)	SD (sec.)
localhost	3593	705
EC2	408	199

From this, we saw an interesting effect of spatial locality. We would have expected the `localhost` calls to be quicker. However, the fact that the Oracle database and EC2 instance are both on us-west, combined with the simple nature of the EC2 server being beefy, made the cloud calls much quicker. This surprised us, but it makes sense in retrospect.

5.2. **CPU versus IO.** We next attempted to discern whether our latency was CPU-bound or IO-bound. We therefore ran 100 sequential requests to load the profile of username `max`, who had few photos and trips, and the profile of username `rigel`, who had a lot of photos and trips.

3

Description	Mean (sec.)	SD (sec.)
`localhost`, `max`	1017	102
EC2, `max`	283	60.6
`localhost`, `rigel`	3593	705
EC2, `rigel`	408	199

From this, we can infer that a lot of latency comes from the Oracle database. Loading all 42 of `rigel`'s photos added a quarter of a second to the page load time on average. This is also likely the reason that `rigel` had a higher deviation on page load times.

5.3. **Sequential versus concurrent.** *Note*: This section is for extra credit, as delineated by the writeup.

Let us explore the time needed to load `rigel`'s profile as a function of concurrent requests.

Concurrency	Min (sec.)	Mean (sec.)	SD (sec.)	Median (sec.)	Max (sec.)
1	311	408	199	338	1843
2	312	427	201	372	2075
3	316	500	233	429	2600
5	327	688	492	578	6925
10	351	744	313	672	3586
15	416	1154	626	968	6492
25	555	1773	806	1601	9173

Methodology: We made requests like `ab -n 1000 -c 10 [url]`, where we kept the number of total requests ten times the number of concurrent requests.

Our application can handle concurrency reasonably well; it handles up to 10 simultaneously heavy requests with reasonable latency. Node's event-driven nature facilitates this. We conjecture that Oracle is not able to handle more than ten simultaneously requests gracefully; we do not have hard evidence of this, but judging by our server's logs, Oracle started clumping together similar database calls at this point. For example, the logs would have seven or eight consecutive responses from the `Photo` table, even though we also made subsequent calls to the `Trip` table (among others).

We finally noticed the long tail on these statistics. Our application will every now and then take a preposterous amount of time to load a page. We aren't positive why this is, but speculate that it may be because of quirks in the OS scheduler on the server or database.

6. Potential Future Extensions

While this web application has a mutltitude of features, there is a lot of room to expand. For example, one feature that is missing from our implementation that is common on most social networks is notifications. Users could be notified of events that happen that are

directly related to them, such as a new destination or photo being added to a trip they are on. Trip requests and friend requests, which are currently rendered in different places, could be combined into this stream.

Another simple feature to extend the application would be sharable links, which would be similar to photos and videos. Users would be able to share links to web pages such as articles, which would then be likeable, rateable, commentable, and would show up on other users newsfeeds. Additionally, users could associate the links with trips or with destinations, the same way albums and destinations are associated with trips.

These are just two basic examples of possible extensions. Of course, the full range of possible extensions is enormous - we could integrate chat, Yelp, hotel booking websites such as Expedia and Kayak, budget management, and so on.